

Single Source Shortest Paths with Negative Weights

Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Mathematics for Data Science 1
Week 12

Dijkstra's algorithm

- Recall the burning pipeline analogy

Dijkstra's algorithm

- Recall the burning pipeline analogy
- We keep track of the following
 - The vertices that have been burnt
 - The expected burn time of vertices

Dijkstra's algorithm

- Recall the burning pipeline analogy
- We keep track of the following
 - The vertices that have been burnt
 - The expected burn time of vertices
- Initially
 - No vertex is burnt
 - Expected burn time of source vertex is 0
 - Expected burn time of rest is ∞

Initialization (assume source vertex 0)

- $B(i) = \text{False}$, for $0 \leq i < n$
 - $UB = \{k \mid B(k) = \text{False}\}$
- $EBT(i) = \begin{cases} 0, & \text{if } i = 0 \\ \infty, & \text{otherwise} \end{cases}$

Dijkstra's algorithm

- Recall the burning pipeline analogy
- We keep track of the following
 - The vertices that have been burnt
 - The expected burn time of vertices
- Initially
 - No vertex is burnt
 - Expected burn time of source vertex is 0
 - Expected burn time of rest is ∞
- While there are vertices yet to burn
 - Pick unburnt vertex with minimum expected burn time, mark it as burnt
 - Update the expected burn time of its neighbours

Initialization (assume source vertex 0)


- $B(i) = \text{False}$, for $0 \leq i < n$
 - $UB = \{k \mid B(k) = \text{False}\}$
- $EBT(i) = \begin{cases} 0, & \text{if } i = 0 \\ \infty, & \text{otherwise} \end{cases}$

Update, if $UB \neq \emptyset$

- Let $j \in UB$ such that $EBT(j) \leq EBT(k)$ for all $k \in UB$
- Update $B(j) = \text{True}$, $UB = UB \setminus \{j\}$
- For each $(j, k) \in E$ such that $k \in UB$,
 $EBT(k) = \min(EBT(k), EBT(j) + W(j, k))$


Correctness requires non-negative edge weights

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt




Correctness requires non-negative edge weights

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is \mathbf{v} , via \mathbf{x}
- Cannot find a shorter path later from \mathbf{y} to \mathbf{v} via \mathbf{w}
 - Burn time of $\mathbf{w} \geq$ burn time of \mathbf{v}
 - Edge from \mathbf{w} to \mathbf{v} has weight ≥ 0




Correctness requires non-negative edge weights

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is \mathbf{v} , via \mathbf{x}
- Cannot find a shorter path later from \mathbf{y} to \mathbf{v} via \mathbf{w}
 - Burn time of $\mathbf{w} \geq$ burn time of \mathbf{v}
 - Edge from \mathbf{w} to \mathbf{v} has weight ≥ 0
- This argument breaks down if edge (\mathbf{w}, \mathbf{v}) can have negative weight




Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt




Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?




Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?
- Recall, negative edge weights are allowed, but no negative cycles




Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?
- Recall, negative edge weights are allowed, but no negative cycles
- Going around a cycle can only add to the length



Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?
- Recall, negative edge weights are allowed, but no negative cycles
- Going around a cycle can only add to the length
- Shortest route to every vertex is a path, no loops



Extending to negative edge weights

- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \cdots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_{\ell}} k$$

- Need not be minimum in terms of number of edges

Extending to negative edge weights

- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_{\ell}} k$$

- Need not be minimum in terms of number of edges
- Every prefix of this path must itself be a minimum weight path
 - $0 \xrightarrow{w_1} j_1$
 - $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2$
 - \dots
 - $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1}$

Extending to negative edge weights

- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_{\ell}} k$$

- Need not be minimum in terms of number of edges
- Every prefix of this path must itself be a minimum weight path

- $0 \xrightarrow{w_1} j_1$

- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2$

- \dots

- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1}$

- Once we discover shortest path to $j_{\ell-1}$, next update will fix shortest path to k

Extending to negative edge weights

- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \cdots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_{\ell}} k$$

- Need not be minimum in terms of number of edges
- Every prefix of this path must itself be a minimum weight path

- $0 \xrightarrow{w_1} j_1$

- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2$

- ...

- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \cdots \xrightarrow{w_{\ell-1}} j_{\ell-1}$

- Once we discover shortest path to $j_{\ell-1}$, next update will fix shortest path to k
- Repeatedly update shortest distance to each vertex based on shortest distance to its neighbours
 - Update cannot push this distance below actual shortest distance

Extending to negative edge weights

- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \cdots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_\ell} k$$

- Need not be minimum in terms of number of edges
- Every prefix of this path must itself be a minimum weight path

- $0 \xrightarrow{w_1} j_1$
- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2$
- \dots
- $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \cdots \xrightarrow{w_{\ell-1}} j_{\ell-1}$

- Once we discover shortest path to $j_{\ell-1}$, next update will fix shortest path to k
- Repeatedly update shortest distance to each vertex based on shortest distance to its neighbours
 - Update cannot push this distance below actual shortest distance
- After ℓ updates, all shortest paths using $\leq \ell$ edges have stabilized
 - Minimum weight path to any node has at most $n-1$ edges
 - After $n-1$ updates, all shortest paths have stabilized

Bellman-Ford Algorithm

Initialization (source vertex 0)

- $D(j)$: minimum distance known so far to vertex j
- $D(j) = \begin{cases} 0, & \text{if } j = 0 \\ \infty, & \text{otherwise} \end{cases}$

Bellman-Ford Algorithm

Initialization (source vertex 0)

- $D(j)$: minimum distance known so far to vertex j
- $D(j) = \begin{cases} 0, & \text{if } j = 0 \\ \infty, & \text{otherwise} \end{cases}$

Repeat $n-1$ times

- For each vertex $j \in \{0, 1, \dots, n-1\}$,
for each edge $(j, k) \in E$,
 $D(k) = \min(D(k),$
 $D(j) + W(j, k))$

Bellman-Ford Algorithm

Initialization (source vertex 0)

- $D(j)$: minimum distance known so far to vertex j
- $D(j) = \begin{cases} 0, & \text{if } j = 0 \\ \infty, & \text{otherwise} \end{cases}$

Repeat $n-1$ times

- For each vertex $j \in \{0, 1, \dots, n-1\}$,
for each edge $(j, k) \in E$,
 $D(k) = \min(D(k),$
 $D(j) + W(j, k))$

Works for directed and undirected graphs


Bellman-Ford Algorithm

Initialization (source vertex 0)

- $D(j)$: minimum distance known so far to vertex j
- $D(j) = \begin{cases} 0, & \text{if } j = 0 \\ \infty, & \text{otherwise} \end{cases}$

Repeat $n-1$ times


- For each vertex $j \in \{0, 1, \dots, n-1\}$,
for each edge $(j, k) \in E$,
 $D(k) = \min(D(k), D(j) + W(j, k))$



Works for directed and undirected graphs

Bellman-Ford Algorithm


v	$D(v)$							
0								
1								
2								
3								
4								
5								
6								
7								



Bellman-Ford Algorithm


v	$D(v)$							
0	0							
1	∞							
2	∞							
3	∞							
4	∞							
5	∞							
6	∞							
7	∞							

- Initialize $D(0) = 0$



Bellman-Ford Algorithm


v	$D(v)$							
0	0	0						
1	∞	10						
2	∞	∞						
3	∞	∞						
4	∞	∞						
5	∞	∞						
6	∞	∞						
7	∞	8						



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0					
1	∞	10	10					
2	∞	∞	∞					
3	∞	∞	∞					
4	∞	∞	∞					
5	∞	∞	12					
6	∞	∞	9					
7	∞	8	8					



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0				
1	∞	10	10	5				
2	∞	∞	∞	10				
3	∞	∞	∞	∞				
4	∞	∞	∞	∞				
5	∞	∞	12	8				
6	∞	∞	9	9				
7	∞	8	8	8				



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0	0			
1	∞	10	10	5	5			
2	∞	∞	∞	10	6			
3	∞	∞	∞	∞	11			
4	∞	∞	∞	∞	∞			
5	∞	∞	12	8	7			
6	∞	∞	9	9	9			
7	∞	8	8	8	8			



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$						
0	0	0	0	0	0	0	
1	∞	10	10	5	5	5	
2	∞	∞	∞	10	6	5	
3	∞	∞	∞	∞	11	7	
4	∞	∞	∞	∞	∞	14	
5	∞	∞	12	8	7	7	
6	∞	∞	9	9	9	9	
7	∞	8	8	8	8	8	



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0	0	0	0	
1	∞	10	10	5	5	5	5	
2	∞	∞	∞	10	6	5	5	
3	∞	∞	∞	∞	11	7	6	
4	∞	∞	∞	∞	∞	14	10	
5	∞	∞	12	8	7	7	7	
6	∞	∞	9	9	9	9	9	
7	∞	8	8	8	8	8	8	



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update
$$D(k) = \min(D(k), D(j) + W(j, k))$$

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



- What if there was a negative cycle?

Bellman-Ford Algorithm


v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



- What if there was a negative cycle?
- Distance would continue to decrease

Bellman-Ford Algorithm

v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



- What if there was a negative cycle?
- Distance would continue to decrease
- Check if update n reduces any $D(v)$

Summary

- Dijkstra's algorithm assumes non-negative edge weights
 - Final distance is frozen each time a vertex “burns”
 - Should not encounter a shorter route discovered later
- Without negative cycles, every shortest route is a path
- Every prefix of a shortest path is also a shortest path
- Iteratively find shortest paths of length $1, 2, \dots, n-1$
- Update distance to each vertex with every iteration — Bellman-Ford algorithm
- If Bellman-Ford algorithm does not converge after $n - 1$ iterations, there is a negative cycle